

Retail-Analysis-with-Walmart-Data

July 1, 2023

```
[2]: # Import necessary libraries
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import dates
from datetime import datetime
```

```
[3]: df = pd.read_csv('Walmart_Store_sales.csv')
df
```

```
[3]:      Store      Date  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price  \
0         1  05-02-2010   1643690.90             0         42.31         2.572
1         1  12-02-2010   1641957.44             1         38.51         2.548
2         1  19-02-2010   1611968.17             0         39.93         2.514
3         1  26-02-2010   1409727.59             0         46.63         2.561
4         1  05-03-2010   1554806.68             0         46.50         2.625
...     ...      ...      ...      ...      ...
6430      45  28-09-2012    713173.95             0         64.88         3.997
6431      45  05-10-2012    733455.07             0         64.89         3.985
6432      45  12-10-2012    734464.36             0         54.47         4.000
6433      45  19-10-2012    718125.53             0         56.47         3.969
6434      45  26-10-2012    760281.43             0         58.85         3.882
```

```
      CPI  Unemployment
0    211.096358      8.106
1    211.242170      8.106
2    211.289143      8.106
3    211.319643      8.106
4    211.350143      8.106
...     ...      ...
6430  192.013558      8.684
6431  192.170412      8.667
6432  192.327265      8.667
6433  192.330854      8.667
6434  192.308899      8.667
```

[6435 rows x 8 columns]

```
[4]: # Convert date to datetime format and show dataset information
df['Date'] = pd.to_datetime(df['Date'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store            6435 non-null   int64
1   Date             6435 non-null   datetime64[ns]
2   Weekly_Sales     6435 non-null   float64
3   Holiday_Flag     6435 non-null   int64
4   Temperature      6435 non-null   float64
5   Fuel_Price       6435 non-null   float64
6   CPI              6435 non-null   float64
7   Unemployment     6435 non-null   float64
dtypes: datetime64[ns](1), float64(5), int64(2)
memory usage: 402.3 KB
```

```
[5]: df.head(2)
```

```
[5]:   Store      Date  Weekly_Sales  Holiday_Flag  Temperature  Fuel_Price  \
0      1  2010-05-02    1643690.90             0         42.31         2.572
1      1  2010-12-02    1641957.44             1         38.51         2.548

      CPI  Unemployment
0  211.096358         8.106
1  211.242170         8.106
```

```
[6]: # checking for missing values
df.isnull().sum()
```

```
[6]: Store            0
Date              0
Weekly_Sales      0
Holiday_Flag      0
Temperature        0
Fuel_Price        0
CPI               0
Unemployment      0
dtype: int64
```

```
[7]: # Splitting Date and create new columns (Day, Month, and Year)
df["Day"] = pd.DatetimeIndex(df['Date']).day
```

```
df['Month'] = pd.DatetimeIndex(df['Date']).month
df['Year'] = pd.DatetimeIndex(df['Date']).year
df
```

```
[7]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
0	1	2010-05-02	1643690.90	0	42.31	2.572	
1	1	2010-12-02	1641957.44	1	38.51	2.548	
2	1	2010-02-19	1611968.17	0	39.93	2.514	
3	1	2010-02-26	1409727.59	0	46.63	2.561	
4	1	2010-05-03	1554806.68	0	46.50	2.625	
...	
6430	45	2012-09-28	713173.95	0	64.88	3.997	
6431	45	2012-05-10	733455.07	0	64.89	3.985	
6432	45	2012-12-10	734464.36	0	54.47	4.000	
6433	45	2012-10-19	718125.53	0	56.47	3.969	
6434	45	2012-10-26	760281.43	0	58.85	3.882	

	CPI	Unemployment	Day	Month	Year
0	211.096358	8.106	2	5	2010
1	211.242170	8.106	2	12	2010
2	211.289143	8.106	19	2	2010
3	211.319643	8.106	26	2	2010
4	211.350143	8.106	3	5	2010
...
6430	192.013558	8.684	28	9	2012
6431	192.170412	8.667	10	5	2012
6432	192.327265	8.667	10	12	2012
6433	192.330854	8.667	19	10	2012
6434	192.308899	8.667	26	10	2012

[6435 rows x 11 columns]

```
[8]: store_sales = df.groupby("Store")["Weekly_Sales"].sum()
store_with_max_sales = store_sales.idxmax()
max_sales = store_sales.max()
print(" Store with max sales:",max_sales )
```

Store with max sales: 301397792.46000004

```
[9]: store_with_min_sales = store_sales.idxmin()
min_sales = store_sales.min()
print(" Store with min sales:",min_sales )
```

Store with min sales: 37160221.960000016

```
[10]: plt.figure(figsize=(15,7))
```

```

# Sum Weekly_Sales for each store, then sortded by total sales
total_sales_for_each_store = df.groupby('Store')['Weekly_Sales'].sum().
    ↪sort_values()
total_sales_for_each_store_array = np.array(total_sales_for_each_store) #
    ↪convert to array

```

<Figure size 1080x504 with 0 Axes>

Which store has minimum and maximum sales?

```

[11]: # Assigning a specific color for the stores have the lowest and highest sales
clrs = ['lightsteelblue' if ((x < max(total_sales_for_each_store_array)) and (x
    ↪ min(total_sales_for_each_store_array))) else 'midnightblue' for x in
    ↪total_sales_for_each_store_array]

ax = total_sales_for_each_store.plot(kind='bar',color=clrs);

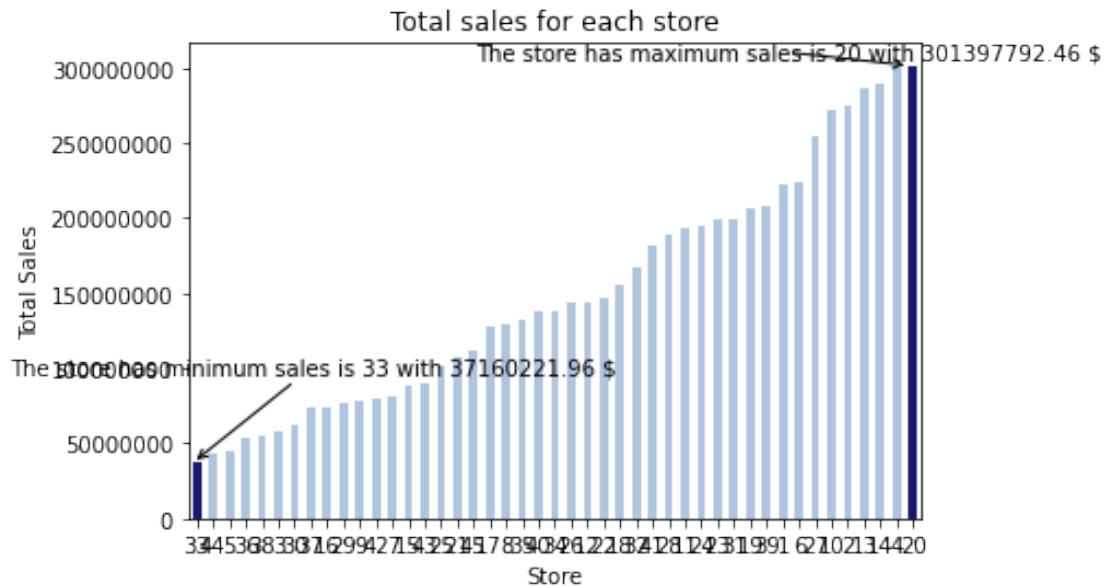
# store have minimum sales
p = ax.patches[0]
print(type(p.get_height()))
ax.annotate("The store has minimum sales is 33 with {0:.2f} $".format((p.
    ↪get_height()), xy=(p.get_x(), p.get_height()), xycoords='data',
    xytext=(0.17, 0.32), textcoords='axes fraction',
    arrowprops=dict(arrowstyle="->", connectionstyle="arc3"),
    horizontalalignment='center', verticalalignment='center')

# store have maximum sales
p = ax.patches[44]
ax.annotate("The store has maximum sales is 20 with {0:.2f} $".format((p.
    ↪get_height()), xy=(p.get_x(), p.get_height()), xycoords='data',
    xytext=(0.82, 0.98), textcoords='axes fraction',
    arrowprops=dict(arrowstyle="->", connectionstyle="arc3"),
    horizontalalignment='center', verticalalignment='center')

# plot properties
plt.xticks(rotation=0)
plt.ticklabel_format(useOffset=False, style='plain', axis='y')
plt.title('Total sales for each store')
plt.xlabel('Store')
plt.ylabel('Total Sales');

```

<class 'numpy.float64'>



Which store has maximum standard deviation i.e., the sales vary a lot. Also, find out the coefficient of mean to standard deviation?

```
[12]: store_std = df.groupby("Store")["Weekly_Sales"].std()
      store_with_max_std = store_std.idxmax()
      max_std = store_std.max()
      print("Store with maximum standard deviation:", store_with_max_std)
```

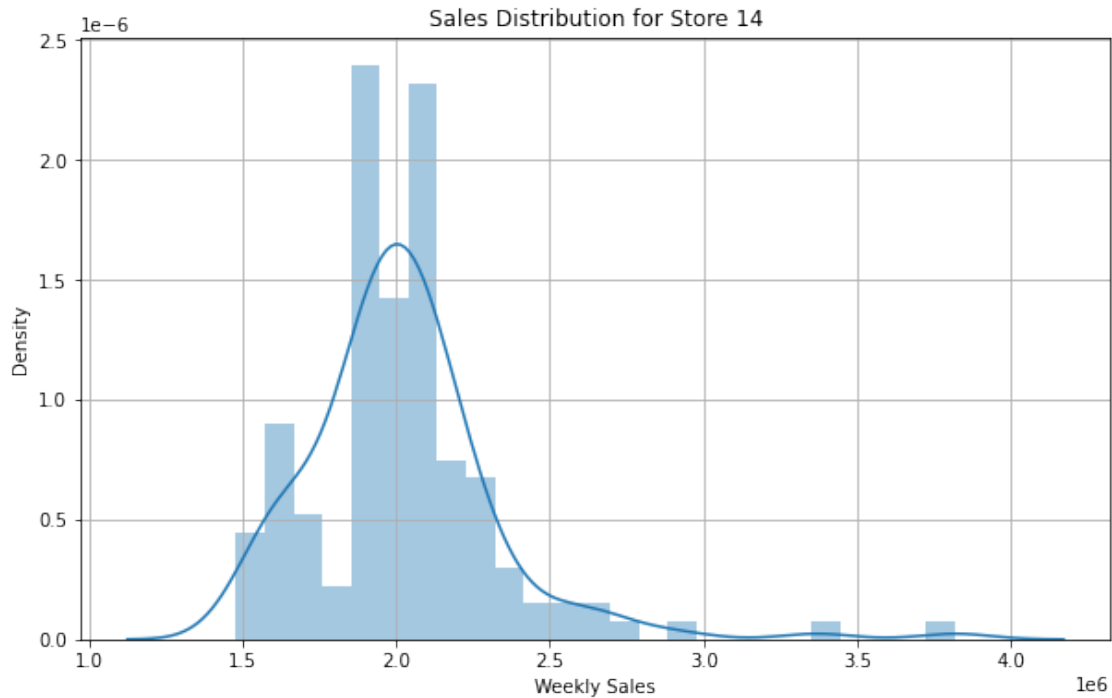
Store with maximum standard deviation: 14

```
[13]: store_data = df[df["Store"] == store_with_max_std]
```

```
[14]: plt.figure(figsize=(10, 6))
      sns.distplot(store_data["Weekly_Sales"])
      plt.title("Sales Distribution for Store {}".format(store_with_max_std))
      plt.xlabel("Weekly Sales")
      plt.ylabel("Density")
      plt.grid(True)
      plt.show()
```

```
/usr/local/lib/python3.7/site-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
```

```
warnings.warn(msg, FutureWarning)
```



```
[15]: # Coefficient of mean to standard deviation
coef_mean_std = pd.DataFrame(df.groupby('Store')['Weekly_Sales'].std() / df.
    ↳groupby('Store')['Weekly_Sales'].mean())
coef_mean_std = coef_mean_std.rename(columns={'Weekly_Sales': 'Coefficient of_
    ↳mean to standard deviation'})
coef_mean_std
```

```
[15]:      Coefficient of mean to standard deviation
Store
1      0.100292
2      0.123424
3      0.115021
4      0.127083
5      0.118668
6      0.135823
7      0.197305
8      0.116953
9      0.126895
10     0.159133
11     0.122262
12     0.137925
13     0.132514
14     0.157137
15     0.193384
```

16	0.165181
17	0.125521
18	0.162845
19	0.132680
20	0.130903
21	0.170292
22	0.156783
23	0.179721
24	0.123637
25	0.159860
26	0.110111
27	0.135155
28	0.137330
29	0.183742
30	0.052008
31	0.090161
32	0.118310
33	0.092868
34	0.108225
35	0.229681
36	0.162579
37	0.042084
38	0.110875
39	0.149908
40	0.123430
41	0.148177
42	0.090335
43	0.064104
44	0.081793
45	0.165613

```
[16]: # Group the data by store and calculate mean and standard deviation of sales
store_stats = df.groupby("Store")["Weekly_Sales"].agg(["mean", "std"])
```

```
[17]: # Calculate the coefficient of mean to standard deviation (coefficient of
      ↪variation)
store_stats["coefficient_of_variation"] = store_stats["mean"] /
      ↪store_stats["std"]
```

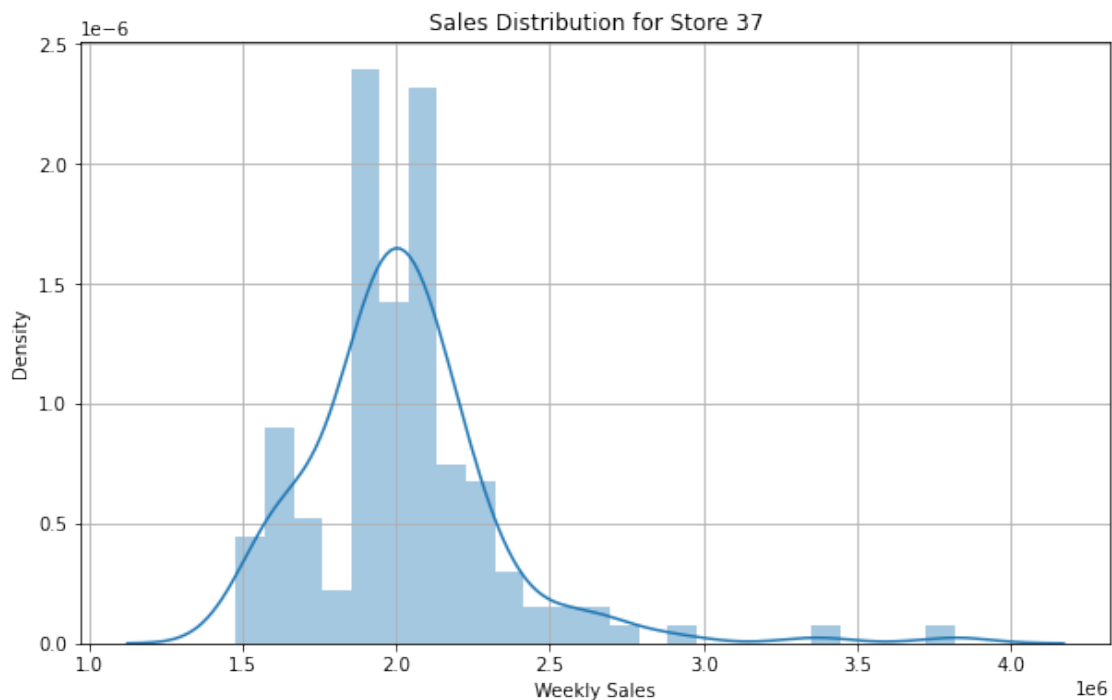
```
[18]: # Find the store with the maximum coefficient of variation
store_with_max_cv = store_stats["coefficient_of_variation"].idxmax()
max_cv = store_stats.loc[store_with_max_cv, "coefficient_of_variation"]
```

```
[19]: print("Maximum coefficient of variation:", max_cv)
```

Maximum coefficient of variation: 23.761932646021123

```
[20]: plt.figure(figsize=(10, 6))
sns.distplot(store_data["Weekly_Sales"])
plt.title("Sales Distribution for Store {}".format(store_with_max_cv))
plt.xlabel("Weekly Sales")
plt.ylabel("Density")
plt.grid(True)
plt.show()
```

/usr/local/lib/python3.7/site-packages/seaborn/distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



Which store/s has good quarterly growth rate in Q3'2012

```
[21]: # Filter the sales data for Q3'2012
q3_2012_data = df[(df["Date"].dt.year == 2012) & (df["Date"].dt.quarter == 3)]
```

```
[22]: # Group the data by store and calculate the quarterly growth rate
store_growth = q3_2012_data.groupby("Store")["Weekly_Sales"].apply(lambda x: (x.
    →iloc[-1] - x.iloc[0]) / x.iloc[0] * 100)
```

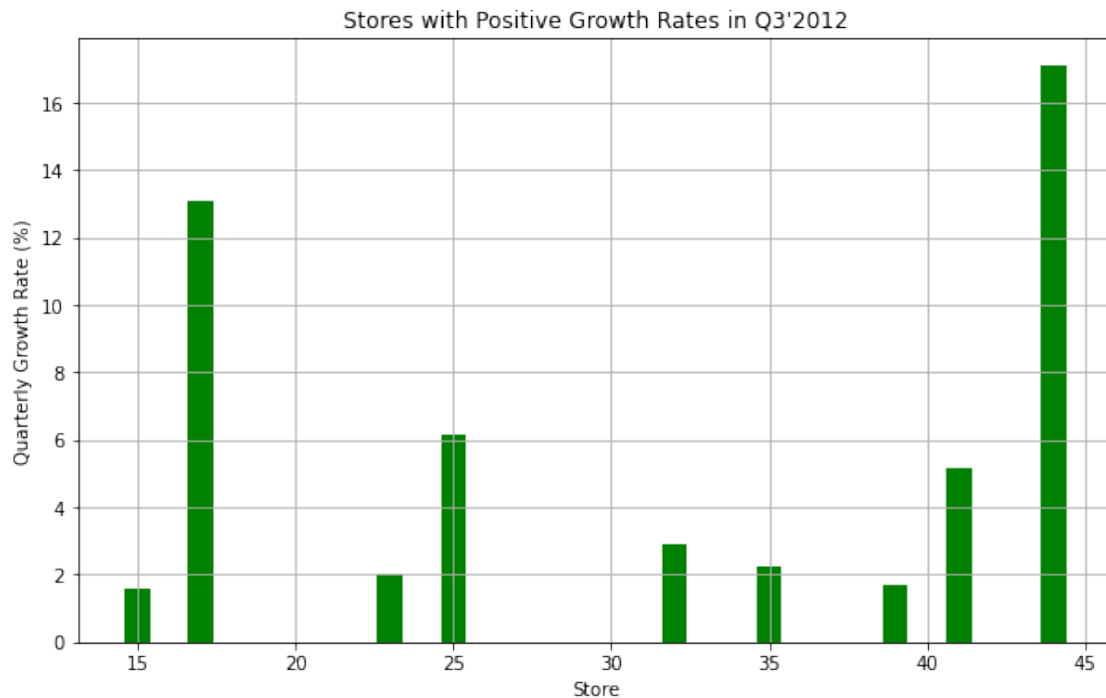


```
[23]: # Find stores with positive quarterly growth rate
stores_with_growth = store_growth[store_growth > 0]
```

```
[24]: print(stores_with_growth)
```

```
Store
15    1.610891
17   13.076958
23    2.035248
25    6.150163
32    2.896583
35    2.228729
39    1.689103
41    5.154565
44   17.093989
Name: Weekly_Sales, dtype: float64
```

```
[25]: plt.figure(figsize=(10, 6))
plt.bar(stores_with_growth.index, stores_with_growth.values, color='green')
plt.title("Stores with Positive Growth Rates in Q3'2012")
plt.xlabel("Store")
plt.ylabel("Quarterly Growth Rate (%)")
plt.grid(True)
plt.show()
```



Find out holidays which have higher sales than the mean sales in non-holiday season for all stores together

```
[27]: # Filter the data for non-holiday weeks and calculate the mean sales
non_holiday_mean = df[df["Holiday_Flag"] == 0]["Weekly_Sales"].mean()

[30]: # Filter the data for holiday weeks and calculate the mean sales for each
      ↪ holiday
holiday_means = df[df["Holiday_Flag"] == 1].groupby("Date")["Weekly_Sales"].
      ↪ mean()

[31]: # Find holidays with higher sales than the mean sales in non-holiday season
holidays_higher_than_mean = holiday_means[holiday_means > non_holiday_mean]

[32]: print(holidays_higher_than_mean)
```

```
Date
2010-11-26    1.462689e+06
2010-12-02    1.074148e+06
2011-11-02    1.051915e+06
2011-11-25    1.479858e+06
2012-07-09    1.074001e+06
2012-10-02    1.111320e+06
Name: Weekly_Sales, dtype: float64
```

```
[33]: # Filter the sales data for each holiday type
super_bowl_data = df[df["Holiday_Flag"] & (df["Date"].isin(["2010-02-12",
      ↪ "2011-02-11", "2012-02-10"]))]
labour_day_data = df[df["Holiday_Flag"] & (df["Date"].isin(["2010-09-10",
      ↪ "2011-09-09", "2012-09-07"]))]
thanksgiving_data = df[df["Holiday_Flag"] & (df["Date"].isin(["2010-11-26",
      ↪ "2011-11-25", "2012-11-23"]))]
christmas_data = df[df["Holiday_Flag"] & (df["Date"].isin(["2010-12-31",
      ↪ "2011-12-30", "2012-12-28"]))]

[34]: # Calculate the mean sales for each holiday type
super_bowl_mean = super_bowl_data.groupby("Date")["Weekly_Sales"].mean()
labour_day_mean = labour_day_data.groupby("Date")["Weekly_Sales"].mean()
thanksgiving_mean = thanksgiving_data.groupby("Date")["Weekly_Sales"].mean()
christmas_mean = christmas_data.groupby("Date")["Weekly_Sales"].mean()

[42]: Super_Bowl_df = pd.DataFrame(df.loc[df.Date.isin(Super_Bowl)]).
      ↪ groupby('Year')['Weekly_Sales'].sum()
Thanksgiving_df = pd.DataFrame(df.loc[df.Date.isin(Thanksgiving)]).
      ↪ groupby('Year')['Weekly_Sales'].sum()
Labour_Day_df = pd.DataFrame(df.loc[df.Date.isin(Labour_Day)]).
      ↪ groupby('Year')['Weekly_Sales'].sum()
```

```

Christmas_df = pd.DataFrame(df.loc[df.Date.isin(Christmas)].
    ↳groupby('Year')['Weekly_Sales'].sum())

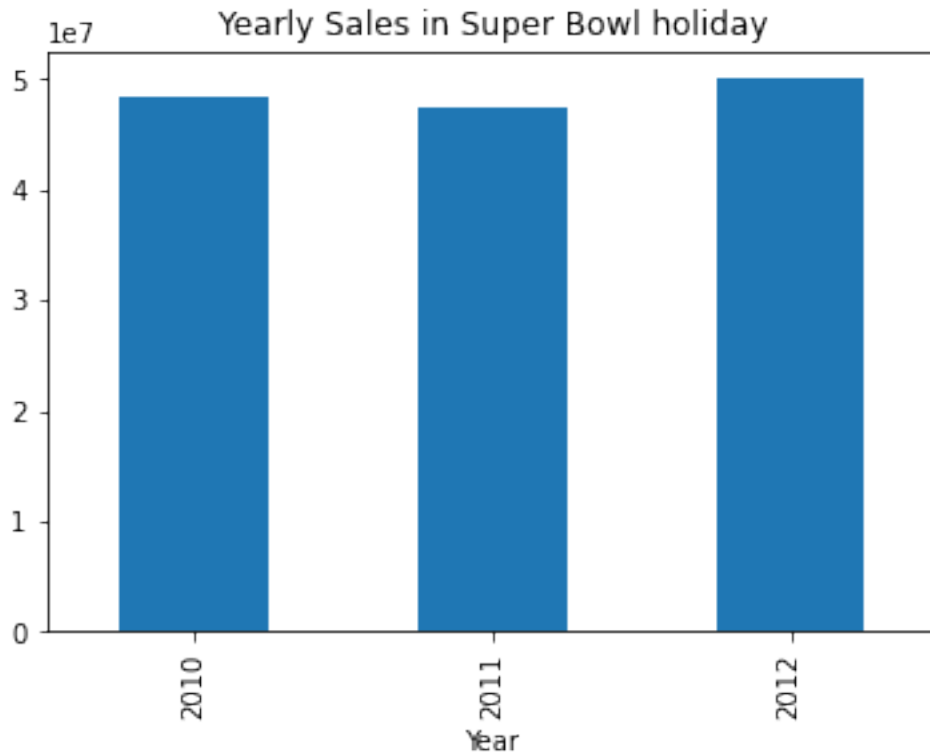
Super_Bowl_df.plot(kind='bar',legend=False,title='Yearly Sales in Super Bowl_
    ↳holiday')
Thanksgiving_df.plot(kind='bar',legend=False,title='Yearly Sales in_
    ↳Thanksgiving holiday')
Labour_Day_df.plot(kind='bar',legend=False,title='Yearly Sales in Labour_Day_
    ↳holiday')
Christmas_df.plot(kind='bar',legend=False,title='Yearly Sales in Christmas_
    ↳holiday')



```


```

[42]: <AxesSubplot:title={'center':'Yearly Sales in Christmas holiday'},
      xlabel='Year'>


```







```
[35]: # Plot the sales trend for each holiday type
plt.figure(figsize=(10, 6))
plt.plot(super_bowl_mean.index, super_bowl_mean.values, label="Super Bowl")
plt.plot(labour_day_mean.index, labour_day_mean.values, label="Labour Day")
plt.plot(thanksgiving_mean.index, thanksgiving_mean.values,
         ↪label="Thanksgiving")
plt.plot(christmas_mean.index, christmas_mean.values, label="Christmas")
plt.title("Sales Trend on Each Holiday Type")
plt.xlabel("Date")
plt.ylabel("Mean Weekly Sales")
plt.legend()
plt.grid(True)
plt.show()
```






```
[38]: def plot_line(df, holiday_dates, holiday_label):
    fig, ax = plt.subplots(figsize = (15,5))
    ax.plot(df['Date'], df['Weekly_Sales'], label=holiday_label)


    for day in holiday_dates:
        day = datetime.strptime(day, '%d-%m-%Y')
        plt.axvline(x=day, linestyle='--', c='r')

    plt.title(holiday_label)
    x_dates = df['Date'].dt.strftime('%Y-%m-%d').sort_values().unique()
    xfmt = dates.DateFormatter('%d-%m-%y')
    ax.xaxis.set_major_formatter(xfmt)
    ax.xaxis.set_major_locator(dates.DayLocator(1))
    plt.gcf().autofmt_xdate(rotation=90)
    plt.show()
```

```
[39]: total_sales = df.groupby('Date')['Weekly_Sales'].sum().reset_index()
Super_Bowl = ['12-2-2010', '11-2-2011', '10-2-2012']
Labour_Day = ['10-9-2010', '9-9-2011', '7-9-2012']
Thanksgiving = ['26-11-2010', '25-11-2011', '23-11-2012']
Christmas = ['31-12-2010', '30-12-2011', '28-12-2012']
```

```
[40]: plot_line(total_sales,Super_Bowl,'Super Bowl')
      plot_line(total_sales,Labour_Day,'Labour Day')
      plot_line(total_sales,Thanksgiving,'Thanksgiving')
      plot_line(total_sales,Christmas,'Christmas')
```






Q5: Provide a monthly and semester view of sales in units and give insights


```
[45]: df["Date"] = pd.to_datetime(df["Date"])
df.set_index("Date", inplace=True)
# Resample the data on a monthly basis to calculate total sales for each month
monthly_sales = df.resample("M")["Weekly_Sales"].sum()
```

```
[46]: # Resample the data on a semester basis to calculate total sales for each
      ↪ semester
semester_sales = df.resample("6M")["Weekly_Sales"].sum()
```

```
[47]: # Plot the monthly sales data
plt.figure(figsize=(10, 6))
plt.plot(monthly_sales.index, monthly_sales.values)
plt.title("Monthly Sales")
plt.xlabel("Date")
plt.ylabel("Total Sales")
plt.grid(True)
plt.show()
```

```
[48]: # Plot the semester sales data
plt.figure(figsize=(10, 6))
plt.plot(semester_sales.index, semester_sales.values)
plt.title("Semester Sales")
plt.xlabel("Date")
plt.ylabel("Total Sales")
plt.grid(True)
plt.show()
```



```
[49]: # Provide insights based on the sales trends
monthly_sales_mean = monthly_sales.mean()
semester_sales_mean = semester_sales.mean()
```

```
[50]: print("Average monthly sales: ", monthly_sales_mean)
print("Average semester sales: ", semester_sales_mean)
```

Average monthly sales: 187144971.86416662
Average semester sales: 962459855.3014281

Build prediction models to forecast demand (Modeling)

```
[51]: # find outliers
fig, axs = plt.subplots(4, figsize=(6,18))
X = df[['Temperature', 'Fuel_Price', 'CPI', 'Unemployment']]
for i, column in enumerate(X):
    sns.boxplot(df[column], ax=axs[i])
```

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only

valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.





FutureWarning

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[54]: # drop the outliers
df = df[(df['Unemployment']<10) & (df['Unemployment']>4.5) &
↳(df['Temperature']>10)]
df
```

```
[54]:
```

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	\
Date						
2010-05-02	1	1643690.90	0	42.31	2.572	
2010-12-02	1	1641957.44	1	38.51	2.548	
2010-02-19	1	1611968.17	0	39.93	2.514	
2010-02-26	1	1409727.59	0	46.63	2.561	
2010-05-03	1	1554806.68	0	46.50	2.625	
...	
2012-09-28	45	713173.95	0	64.88	3.997	
2012-05-10	45	733455.07	0	64.89	3.985	
2012-12-10	45	734464.36	0	54.47	4.000	
2012-10-19	45	718125.53	0	56.47	3.969	
2012-10-26	45	760281.43	0	58.85	3.882	

	CPI	Unemployment	Day	Month	Year
Date					
2010-05-02	211.096358	8.106	2	5	2010
2010-12-02	211.242170	8.106	2	12	2010
2010-02-19	211.289143	8.106	19	2	2010
2010-02-26	211.319643	8.106	26	2	2010
2010-05-03	211.350143	8.106	3	5	2010
...
2012-09-28	192.013558	8.684	28	9	2012
2012-05-10	192.170412	8.667	10	5	2012
2012-12-10	192.327265	8.667	10	12	2012
2012-10-19	192.330854	8.667	19	10	2012
2012-10-26	192.308899	8.667	26	10	2012

[5658 rows x 10 columns]

```
[56]: # check outliers
fig, axs = plt.subplots(4,figsize=(6,18))
X = df[['Temperature','Fuel_Price','CPI','Unemployment']]
for i,column in enumerate(X):
    sns.boxplot(df[column], ax=axs[i])
```





/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.

FutureWarning
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.

FutureWarning
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variable as a keyword arg: x. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.

FutureWarning



Building model

```
[59]: # Import sklearn
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.linear_model import LinearRegression

[60]: # Select features and target
X = df[['Store', 'Fuel_Price', 'CPI', 'Unemployment', 'Day', 'Month', 'Year']]
y = df['Weekly_Sales']

# Split data to train and test (0.80:0.20)
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2)

[61]: # Linear Regression model
print('Linear Regression:')
print()
reg = LinearRegression()
reg.fit(X_train, y_train)
y_pred = reg.predict(X_test)
print('Accuracy:', reg.score(X_train, y_train)*100)

Linear Regression:

Accuracy: 13.54426145850528

[62]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))

sns.scatterplot(y_pred, y_test);
```


Mean Absolute Error: 453254.98996341514

Mean Squared Error: 302364524014.50336

Root Mean Squared Error: 549876.8262206577

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variables as keyword args: x, y. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[63]: # Random Forest Regressor
print('Random Forest Regressor:')
print()
rfr = RandomForestRegressor(n_estimators = 400,max_depth=15,n_jobs=5)
rfr.fit(X_train,y_train)
y_pred=rfr.predict(X_test)
print('Accuracy:',rfr.score(X_test, y_test)*100)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_pred)))

sns.scatterplot(y_pred, y_test);
```

Random Forest Regressor:

Accuracy: 94.2896058424671

Mean Absolute Error: 70803.25902171675


Mean Squared Error: 19246618028.373486

Root Mean Squared Error: 138732.18094001652

/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variables as keyword args: x, y. From version 0.12, the only

valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



```
[ ]: Random Forest is the best model with accuracy 94%
```